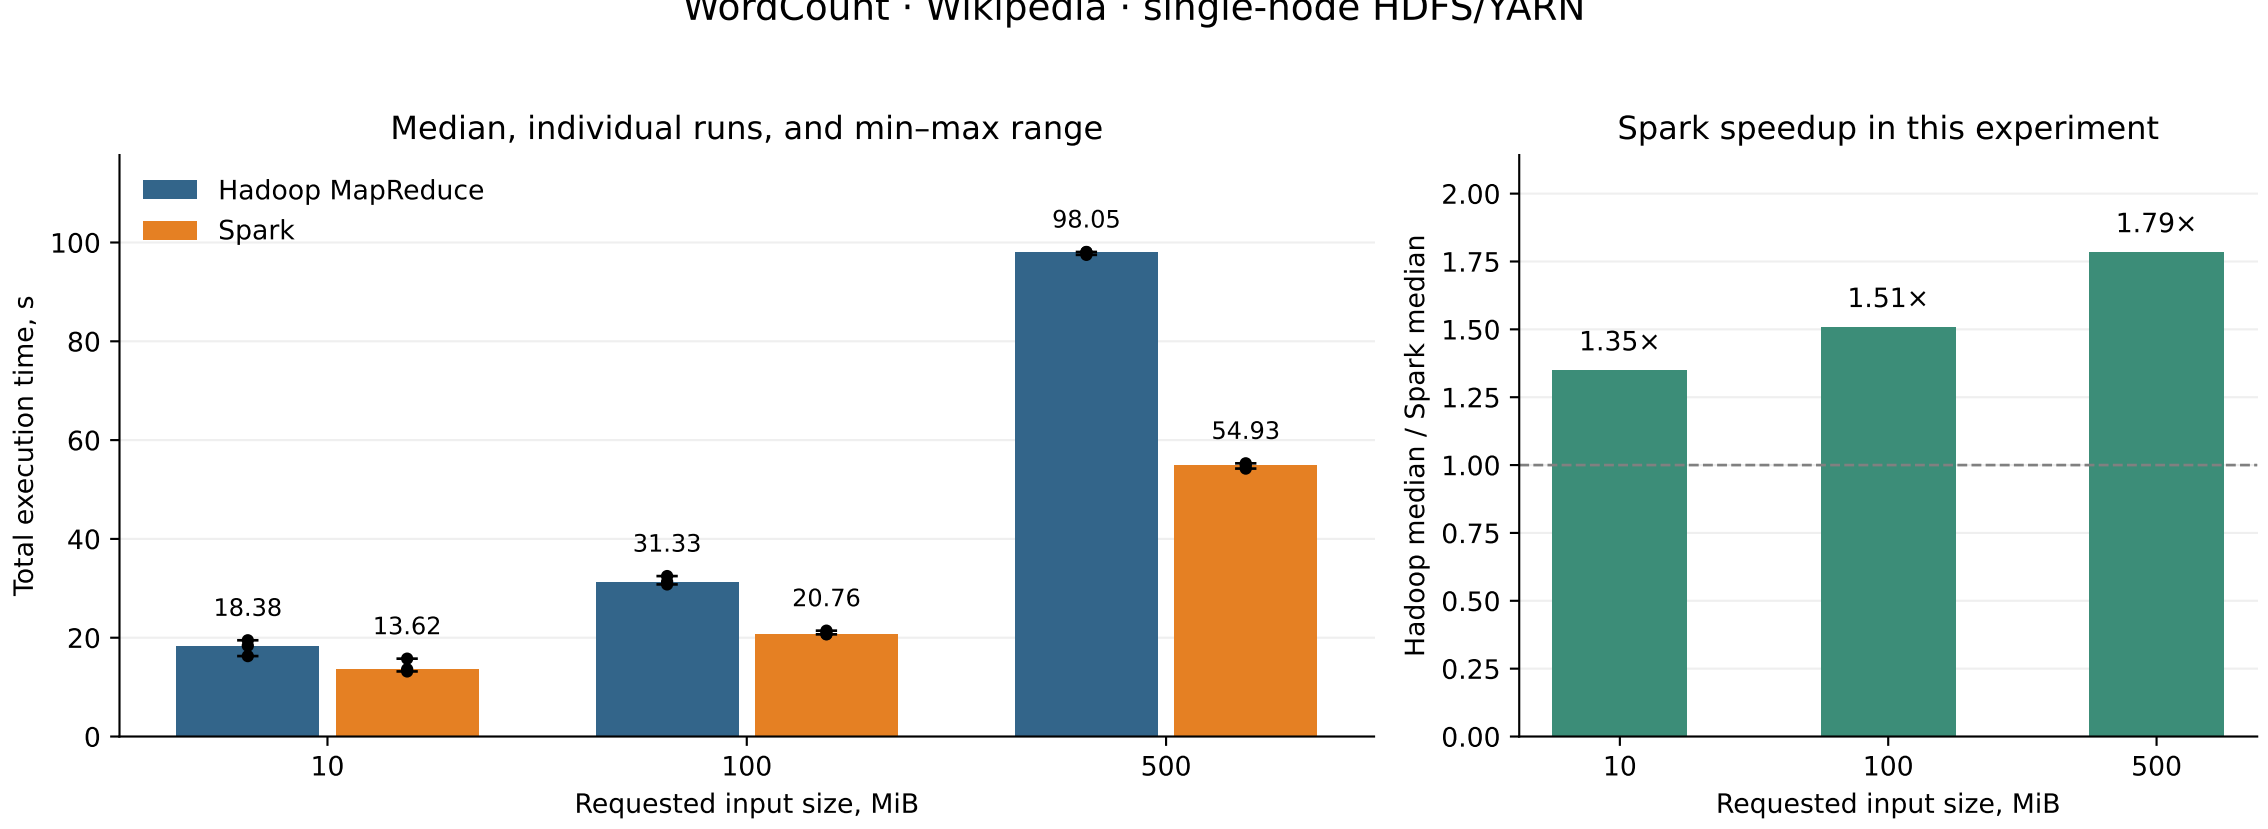
WordCount · Wikipedia · single-node HDFS/YARN
Median, individual runs, and min–max range
Hadoop MapReduce
Spark
Total execution time, s
0
20
40
60
80
100
18.38
13.62
31.33
20.76
98.05
54.93
10
100
500
Requested input size, MiB
Spark speedup in this experiment
Hadoop median / Spark median
0.00
0.25
0.50
0.75
1.00
1.25
1.50
1.75
2.00
1.35×
1.51×
1.79×
10
100
500
Requested input size, MiB
3 repetitions; includes application startup, reading, processing, and output writing.